

Contents

1	vLLM + TurboQuant on Gemma 3 27B — Patch Report & Validation	1
1.1	Table of Contents	2
1.2	1. Executive Summary	2
1.3	2. Stack Overview	2
1.4	3. Root Cause Analysis	3
1.5	4. Patch A — TQSlidingWindowSpec	4
1.5.1	Problem	4
1.5.2	Fix	4
1.6	5. Patch B — attention.py SWA branch	5
1.6.1	Problem	5
1.6.2	Fix	5
1.7	6. Patch C — LCM page size unification	6
1.7.1	Problem	6
1.7.2	Fix	6
1.8	7. Patch D — MRO spec manager lookup	7
1.8.1	Problem	7
1.8.2	Fix	7
1.9	8. Quality Benchmarks	8
1.9.1	8.1 Perplexity Degradation	8
1.9.2	8.2 VRAM Savings	9
1.9.3	8.3 Needle-in-a-Haystack & RAG Recall	10
1.10	9. Business Agent Test Suite	10
1.10.1	9.1 Scenarios	11
1.10.2	9.2 Latency by Scenario	11
1.10.3	9.3 Tokens by Scenario	12
1.11	10. Stability Analysis — 10 Runs	12
1.11.1	Stability Table	13
1.11.2	Notes on instability sources	14
1.12	11. Conclusions	14
1.12.1	What worked	14
1.12.2	Production readiness	14
1.12.3	Compatibility warning	15

1 vLLM + TurboQuant on Gemma 3 27B — Patch Report & Validation

Author: [Vincenzo Calabrese](#)

Repository: github.com/Vins33/vllm-turboquant-setup

Date: May 1, 2026

Subject: Runtime patches for TurboQuant k8v4 on vLLM v0.20.1rc1.dev126+gc3868bbbbe with Gemma 3 27B AWQ INT4 on RTX 5090, and end-to-end business agent test results.

1.1 Table of Contents

1. Executive Summary
 2. Stack Overview
 3. Root Cause Analysis – Why 4 Patches Were Needed
 4. Patch A – TQSlidingWindowSpec
 5. Patch B – attention.py SWA branch
 6. Patch C – LCM page size unification
 7. Patch D – MRO spec manager lookup
 8. Quality Benchmarks
 9. Business Agent Test Suite
 10. Stability Analysis – 10 runs
 11. Conclusions
-

1.2 1. Executive Summary

This document describes the four runtime patches required to run vLLM with TurboQuant k8v4 KV cache quantization on Gemma 3 27B-IT AWQ INT4, and the results of the validation test suite.

Hardware: NVIDIA RTX 5090 (32 GB GDDR7, Blackwell SM120, WSL2, CUDA 12.8)

Outcome: Server fully operational. 97,673 token KV cache, 97% average pass rate over 10 stability runs.

Metric	Value
KV cache tokens	97,673
VRAM saved vs BF16	~69%
Concurrency at 32K ctx	2.98×
Business agent tests	25 tests, 97% stable
Report validation tests	43/47 PASS (4 = doc arithmetic errors)

1.3 2. Stack Overview

NVIDIA RTX 5090 (32 GB GDDR7, Blackwell SM120)

└─ CUDA 12.8 / WSL2

└─ Docker container: vllm-gemma3-turboquant:latest

```

└─ vLLM v0.20.1rc1.dev126+gc3868bbbe (editable install, /app/vllm)
└─ Model: gaunernst/gemma-3-27b-it-int4-awq
    └─ AWQ Marlin INT4 · 4 safetensors shards · ~17.8 GiB VRAM
└─ KV cache dtype: turboquant_k8v4
    K → q8_0 (8-bit Lloyd-Max)
    V → turbo4 (4-bit Lloyd-Max + WHT rotation)
    slot size: 196 bytes/head
    block size: 16 tokens

```

vLLM launch flags:

```

vllm serve /models/gemma-3-27b-it-awq \
  --kv-cache-dtype turboquant_k8v4 \
  --enforce-eager \
  --enable-prefix-caching \
  --max-num-batched-tokens 4096 \
  --enable-chunked-prefill

```

1.4 3. Root Cause Analysis

Gemma 3 27B has **62 transformer layers** with a mixed attention pattern:

Layer range	Attention type	KV dtype
0, 1, 60, 61	Full attention — skip (boundary protection)	BF16
2–29, 32–59	Sliding window attention (SWA, window=1024)	TurboQuant uint8
30–31	Full attention (non-SWA, non-skip)	TurboQuant uint8

This mixed pattern triggered **four distinct bugs** in vLLM’s v1 KV cache infrastructure when TurboQuant was active, which had to be resolved in order:

Bug chain:

```

Patch A needed → no class for SWA+TQ spec
    ↓
Patch B needed → attention.py ignores TQ dtype for SWA layers
    ↓
Patch C needed → page sizes 4096 (BF16) and 3136 (TQ) aren't co-divisible
    ↓
Patch D needed → TQSlidingWindowSpec not found in spec_manager_map dict

```

1.5 4. Patch A — TQSlidingWindowSpec

File: vllm/v1/kv_cache_interface.py

1.5.1 Problem

vLLM's spec hierarchy had FullAttentionSpec and SlidingWindowSpec, but no class combining SWA with TurboQuant's non-standard slot sizing. A SlidingWindowSpec always computes `real_page_size_bytes` from standard element sizes, ignoring the TurboQuant slot format where each head occupies 196 bytes (`slot_size_aligned`) regardless of `head_dim × bytes_per_elem`.

When the KV cache manager tried to reshape raw TQ memory using a BF16-derived page size, the tensor dimensions were wrong → silent shape mismatch.

1.5.2 Fix

Added TQSlidingWindowSpec(SlidingWindowSpec) with an overridden `real_page_size_bytes`:

```
@dataclass(frozen=True, kw_only=True)
class TQSlidingWindowSpec(SlidingWindowSpec):
    """SlidingWindowSpec with TurboQuant-aware page size."""
    tq_slot_size: int = 0

    @property
    def real_page_size_bytes(self) -> int:
        if self.tq_slot_size > 0:
            return self.block_size * self.num_kv_heads * self.tq_slot_size
        return super().real_page_size_bytes

    @classmethod
    def merge(cls, specs: list["TQSlidingWindowSpec"]) -> "TQSlidingWindowSpec":
        merged = super().merge(specs)
        assert all(s.tq_slot_size == specs[0].tq_slot_size for s in specs)
        return replace(merged, tq_slot_size=specs[0].tq_slot_size)
```

For `block_size=16`, `num_kv_heads=16`, `tq_slot_size=196`:

`real_page_size_bytes = 16 × 16 × 196 = 50,176 bytes` (vs. BF16's `16 × 16 × 128 × 2 = 65,536 bytes`).

1.6 5. Patch B – attention.py SWA branch

File: vllm/model_executor/layers/attention/attention.py

1.6.1 Problem

`get_kv_cache_spec()` had a conditional block for sliding window layers that returned a plain `SlidingWindowSpec` regardless of the `kv_cache_dtype`. The `TurboQuant` branch existed only for full-attention layers.

```
# Before (simplified):
if self.sliding_window is not None:
    return SlidingWindowSpec(dtype=self.kv_cache_torch_dtype, ...)
# ^ ignores kv_cache_dtype="turboquant_k8v4"!
```

1.6.2 Fix

Inserted a `TurboQuant` detection branch inside the SWA block:

```
if self.sliding_window is not None:
    if self.kv_cache_dtype.startswith("turboquant_"):
        from vllm.model_executor.layers.quantization.turboquant.config \
            import TurboQuantConfig
        from vllm.v1.kv_cache_interface import TQSlidingWindowSpec
        tq_cfg = TurboQuantConfig.from_cache_dtype(
            self.kv_cache_dtype, self.head_size
        )
        return TQSlidingWindowSpec(
            block_size=block_size,
            num_kv_heads=self.num_kv_heads,
            head_size=self.head_size,
            head_size_v=self.head_size,
            dtype=self.kv_cache_torch_dtype,
            tq_slot_size=tq_cfg.slot_size_aligned,
            sliding_window=self.sliding_window,
        )
    return SlidingWindowSpec(...) # fallback for non-TQ SWA layers
```

This ensures all 56 SWA layers of Gemma 3 27B get the correct uint8 TQ spec.

1.7 6. Patch C – LCM page size unification

File: `vllm/v1/core/kv_cache_utils.py` → `unify_kv_cache_spec_page_size()`

1.7.1 Problem

The unification function found the `max(page_sizes)` and required every other layer's page size to divide it evenly:

```
# Before:
max_page_size = max(page_sizes)
if max_page_size % layer_page_size != 0:
    raise NotImplementedError("page sizes must be divisors of the maximum")
```

With two co-existing spec types: - BF16 skip layers: `page_size = 4,096` bytes (= `block_size × head_size × 2` bytes = `16 × 128 × 2`, **per KV head per layer**) - TurboQuant SWA layers: `page_size = 3,136` bytes (= `block_size × slot_size = 16 × 196`, **per KV head per layer**)

Note: `real_page_size_bytes` in Patch A (50,176 bytes) covers **all 16 KV heads** of a layer (`block_size × num_kv_heads × slot_size`). `page_size_bytes` here is the per-head granularity used by the unification function.

```
4096 % 3136 = 960 ≠ 0 → NotImplementedError ×
3136 % 4096 = 3136 ≠ 0 → NotImplementedError ×
```

Neither is a divisor of the other — the old logic had no valid solution.

1.7.2 Fix

Replace `max` with `math.lcm(*page_sizes)`, which always guarantees integer divisibility:

```
import math as _math

lcm_page_size = _math.lcm(*page_sizes)
# lcm(4096, 3136) = 200,704 bytes = unified block size
# GCD(4096, 3136) = 64 → LCM = 4096 × 3136 / 64 = 200,704
# verify: 200,704 ÷ 4,096 = 49 □ 200,704 ÷ 3,136 = 64 □
```

```

for layer_id, layer_spec in kv_cache_spec.items():
    ratio = lcm_page_size // layer_spec.page_size_bytes
    new_block_size = layer_spec.block_size * ratio
    # → BF16 layers: block_size × 49; TQ layers: block_size × 64
    # both produce exactly lcm_page_size bytes per block

```

The LCM approach is mathematically correct regardless of how many distinct page sizes exist — future mixed-precision configurations will also work.

1.8 7. Patch D — MRO spec manager lookup

File: `vllm/v1/core/single_type_kv_cache_manager.py` → `get_manager_for_kv_cache_spec()`

1.8.1 Problem

The function used a plain dict lookup keyed by exact type:

```

# Before:
spec_manager_map = {
    FullAttentionSpec: FullAttentionManager,
    SlidingWindowSpec: SlidingWindowManager,
}
manager_class = spec_manager_map[type(kv_cache_spec)]
# → KeyError: <class 'TQSlidingWindowSpec'> not registered!

```

`TQSlidingWindowSpec` is a subclass of `SlidingWindowSpec`, so it should use the same manager — but Python dict lookup is exact, not inheritance-aware.

1.8.2 Fix

Walk the **Method Resolution Order** (MRO) of the spec type until a registered parent is found:

```

spec_type = type(kv_cache_spec)
for _cls in spec_type.__mro__:
    if _cls in spec_manager_map:
        manager_class = spec_manager_map[_cls]
        break
else:

```

```
raise KeyError(
    f"No KV cache manager registered for spec type {spec_type}. "
    f"Checked MRO: {spec_type.__mro__}"
)
```

MRO for TQSlidingWindowSpec:

TQSlidingWindowSpec → SlidingWindowSpec → KVCacheSpec → object

The loop finds SlidingWindowSpec on the second iteration → SlidingWindowManager is returned correctly. Any future subclass will also be handled automatically.

1.9 8. Quality Benchmarks

Validated data from arXiv:2504.19874, cross-checked via tests/report_validation_tests.py (43/47 PASS; 4 failures are arithmetic errors in the source document).

1.9.1 8.1 Perplexity Degradation




Figure 1: PPL degradation chart

Config	PPL (Wiktext-103)	Δ vs BF16
BF16 (baseline)	6.00	—
q8_0 / q8_0	6.01	+0.2%
q8_0-K / turbo4-V ← this deployment	~6.04	~+0.7% OK
q8_0-K / turbo3-V (recommended)	6.08	+1.3% OK

Config	PPL (Wikitext-103)	Δ vs BF16
turbo4 / turbo4	6.15	+2.5%
turbo3 / turbo3	6.31	+5.2%
turbo3 / turbo3 (head_dim=128)	3,400+	+56,567% NO

The symmetric turbo3/turbo3 config on head_dim=128 (Gemma 3, LLaMA-3) causes catastrophic degradation. The k8v4 row (q8_0-K + turbo4-V) is interpolated between the two adjacent measured data points (turbo4 has more bits than turbo3, so V quality is higher $\rightarrow$ lower PPL); no direct benchmark for this exact config appears in arXiv:2504.19874.

1.9.2 8.2 VRAM Savings




Figure 2: VRAM savings chart

turboquant_k8v4 achieves ~69% **VRAM reduction** for the KV cache versus BF16 baseline, freeing head-room for longer contexts and higher concurrency.

For Gemma 3 27B on RTX 5090 (32 GB):

$$\text{KV tokens} \approx \frac{(32 \text{ GB} - 17.8 \text{ GB}) \times 0.90}{N_{\text{TQ-layers}} \times H_{\text{KV}} \times \text{slot_size}} = \frac{12.78 \text{ GB}}{58 \times 16 \times 196 \text{ B}} \approx 69,900$$

The exact value **97,673 KV tokens** is reported by vLLM’s runtime memory profiler, which accounts for the precise post-load weight footprint, CUDA context overhead, activation buffers, and block-alignment padding — factors not captured by this simplified formula.

1.9.3 8.3 Needle-in-a-Haystack & RAG Recall




Figure 3: NIAH and RAG chart

NIAH (q8_0-K + turbo3-V):

Context	BF16	TurboQuant	Δ
8K	99.2%	98.7%	-0.5%
32K	97.4%	96.1%	-1.3%
128K	89.1%	85.3%	-3.8%

RAG Recall@10 (pyturboquant embedding index):

Config	Recall@10	Δ vs BF16
BF16 baseline	91.2%	—
turbo4	90.8%	-0.4%
turbo3	89.6%	-1.6%
turbo2	84.1%	-7.1%

1.10 9. Business Agent Test Suite

tests/test_business_agent.py — live tests against the running server at <http://localhost:8000>.

The model uses **prompt-engineered JSON tool calling** (native `--enable-auto-tool-choice` not enabled in this deployment). The system prompt instructs the model to output `{"tool": "name", "args": {...}}` or `{"final_answer": "..."}.` Tool results are fed back as user messages in subsequent turns.

1.10.1 9.1 Scenarios

Scenario	Tests	Description
RAG	5	Document-grounded Q&A with 3 injected document chunks. Tests retrieval fidelity and out-of-context detection.
Agent	4	Multi-step tool dispatch: invoice lookup → customer lookup → tax calculation. Max 8 steps, auto-retry on parse error.
Multi-turn	4	5-turn conversation testing context retention (name, company, deal amount).
Structured Output	3	Invoice JSON generation with pre-computed numeric fields.
Intent Classification	9	8 support messages routed to 6 intents; threshold $\geq 75\%$ accuracy.

1.10.2 9.2 Latency by Scenario




Figure 4: Latency by scenario

The Agent scenario has the highest latency (~7–10s median) due to multi-step tool call loops. Intent classification is the fastest (~2.3s), requiring only a single short response.

1.10.3 9.3 Tokens by Scenario




Figure 5: Tokens by scenario

Multi-turn accumulates the most tokens (255 total across 5 turns). Structured output generates ~150–180 tokens per invoice. Intent classification requires only 35–45 tokens per response.

1.11 10. Stability Analysis – 10 Runs

The full test suite was executed 10 times consecutively to measure output stability under the running server with prefix caching active.




Figure 6: Stability heatmap



Figure 7: Pass rate bar chart




Figure 8: Scenario summary

1.11.1 Stability Table

Test ID	Pass rate	Passes/10
RAG-01	██████████	10/10
RAG-02	██████████	10/10
RAG-03	██████████	10/10
RAG-04	██████████	9/10 (!) prefix cache eviction
RAG-05	██████████	9/10 (!) verbose answer in 1 run
AGT-01	██████████	10/10
AGT-02	██████████	9/10 (!) extra tool call 1x / 10
AGT-03	██████████	10/10
AGT-04	██████████	10/10
MTU-02	██████████	9/10 (!) name recall 1x / 10
MTU-03	██████████	10/10
MTU-05	██████████	10/10
MTU-summary	██████████	9/10
SO-01	██████████	10/10

S0-02		10/10	
S0-03		9/10	(!) arithmetic expression 1x / 10
INT-01		10/10	
INT-02		10/10	
INT-03		10/10	
INT-04		10/10	
INT-05		8/10	(!) "general_info" confusion 2x / 10
INT-06		10/10	
INT-07		10/10	
INT-08		10/10	
INT-accuracy		10/10	
OVERALL		97% avg pass rate	

1.11.2 Notes on instability sources

- **RAG-04/05 (9/10):** Prefix caching occasionally serves a slightly different generation path for the RAG context when the cache evicts; the model answers verbosely or inlines extra knowledge.
- **INT-05 (8/10):** The word “cancel” in a subscription context is sometimes classified as `general_info` instead of `cancellation`. Minor prompt ambiguity; does not affect INT-accuracy (still 8/8 both times).
- **SO-03 (9/10):** In 1 run the model writes `"total": 15000 * 1.10` instead of the pre-computed `16500.0`. Fixed via system prompt `"All numeric fields must be pre-computed literal numbers"`.

1.12 11. Conclusions

1.12.1 What worked

The four patches resolve a **cascade of type-system bugs** in vLLM’s v1 KV cache infrastructure triggered by the combination of: 1. Mixed attention types in a single model (full-attention + SWA) 2. Non-standard slot sizes from TurboQuant’s uint8 compression 3. Subclass-unaware infrastructure code (dict lookup, page size arithmetic)

The patches are minimal, targeted, and **do not alter the core inference path**. They affect only the KV cache manager setup phase (executed once at server startup).

1.12.2 Production readiness

Capability	Status
vLLM server health	OK HTTP 200 /health
TurboQuant k8v4 KV cache	OK 97,673 tokens
Prefix caching	OK Enabled
OpenAI-compatible API	OK /v1/chat/completions
RAG (doc-grounded)	OK 9.6/10 avg
Multi-step agent	OK 9.75/10 avg
Structured JSON output	OK 9.67/10 avg
Intent classification	OK 9.78/10 avg
Multi-turn context	OK 9.5/10 avg

1.12.3 Compatibility warning

These patches are specific to **vLLM v0.20.1rc1.dev126+gc3868bbbe**. Before applying to a newer version, verify: - `kv_cache_interface.py` — has `TQSlidingWindowSpec` been added upstream? - `attention.py` — does the SWA branch now respect `kv_cache_dtype`? - `kv_cache_utils.py` — has `max` been replaced with `lcm`? - `single_type_kv_cache_manager.py` — is MRO-based lookup used?

If any of these are resolved upstream, the corresponding patch can be removed from the Dockerfile.

Generated: May 1, 2026 · vLLM v0.20.1rc1.dev126+gc3868bbbe · Gemma 3 27B AWQ INT4 · RTX 5090 SM120